

PROJECT PITCH

RESOLVIA

Investment Banking Exception Intelligence & Resolution

From Break to Verified Resolution

Business Proposal & Industry Case

Hiral Sarkar (PGA 45)

PG Data Science and Analytics with Gen AI

September 2026

1. Executive Pitch

Investment banks and asset managers have invested heavily in reconciliation, matching and straight-through processing. The business problem has therefore shifted: once a system identifies a break, how quickly and safely can the organisation understand it, resolve it and prove that it is actually fixed?

RESOLVIA proposes an evidence-grounded agentic layer for that gap. It does not replace reconciliation. It sits after break detection and combines structured-data ML, OCR, retrieval, historical precedent, LLM reasoning and workflow controls to investigate exceptions, recommend corrective action and validate closure.

The commercial proposition is simple: automate the investigation work that is repetitive enough for rules but too heterogeneous for conventional RPA; reserve human judgment for consequential decisions; and turn every resolution into a traceable evidence trail.

2. Why This Matters to a Financial Institution

Pressure	Industry evidence	Business consequence
High exception volume	SS&C GlobeOp: 800K trade breaks reviewed/month	Large recurring workload; small per-break savings can create material capacity.
Tight post-trade windows	DTCC says faster automation is critical under T+1	Delayed investigation can increase settlement risk and escalation.
Automation is incomplete	Broadridge: 16% manual, 53% partial, 31% mostly automated, 0% fully automated in its FCM survey	Residual exception queue remains.
Repeatable automation works	CRISIL: 65% independent break resolution in one case	Next opportunity is expanding the addressable exception set.
Capacity is measurable	SS&C: 3,680 hours/month returned; CRISIL: ~80 onshore hours/week saved in another case	Benefits can be translated into capacity and cost avoidance.
Control matters	IVP: rules + agents + human oversight and auditability	Agentic automation needs governance and approval thresholds.

3. The Industry Automation Ladder

Level	What is automated	Evidence	Gap	RESOLVIA
0 - Matching	Compare records, tolerances, identify match/break	DTCC: 98.4% same-day equity matching in Q1 2026; >99% same-day agreement as ITP best practice	Break still requires investigation	Upstream dependency
1 - Manual resolution	Analyst gathers evidence and coordinates correction	Broadridge FCM survey: 16% predominantly manual; 53% with manual investigation	High analyst touch	Controlled baseline
2 - RPA / rules	Known patterns, data movement, comments, deterministic actions	SS&C and CRISIL cases show large-scale post-break automation	Static rules struggle with novel evidence-heavy cases	Measure rules ceiling
3 - Agentic AI	Dynamic retrieval, contextual reasoning, RCA, recommendation, orchestration	IVP and Zensar/AWS describe commercial agentic approaches	Need proof of incremental value and governance	Evaluate safely against Level 2

4. Level 0 - Matching & Reconciliation

Matching technology is already sophisticated. DTCC reports 98.4% same-day matching for equities among its clients in Q1 2026 and 2.6 million CTM transactions processed daily. DTCC also publishes best-practice measures including more than 99% same-day agreement and more than 90% of block exceptions fixed within one hour.

This matters because RESOLVIA should not compete with a problem the industry has already solved well. The opportunity begins after a genuine exception has been created.

5. Level 1 - Manual Break Resolution

Broadridge's State of the Market survey reports 16% of respondents handling trade breaks predominantly manually, 53% with some manual investigation, 31% mostly automated and 0% fully automated. This is specifically an FCM survey and should not be generalised to all investment banks.

The residual workload is investigative: connecting trade records, confirmations, settlement status, custodian statements, cash movements, emails and historical cases before deciding what happened. That cognitive chain is RESOLVIA's target.

6. Level 2 - Conventional Automation / RPA

SS&C GlobeOp

SS&C GlobeOp deployed digital workers to investigate existing trade breaks for 150 clients. SS&C reports 800,000 trade breaks reviewed per month, 3,680 hours returned per month, a 20% increase in day-zero break resolution and 100% accuracy in the reported case. The workers classify existing breaks against predefined rules, standardise comments and escalate complex items.

CRISIL Integral IQ

CRISIL reports a 2025 case with 65% of breaks resolved independently offshore across transactions, cash and positions. Another case reports 90%+ break-resolution accuracy, 50% lower reconciliation time and approximately 55% cost savings.

A separate 2025 case reports a 99% reconciliation TAT reduction, from up to 1,200 seconds per fund to under 10 seconds per fund, plus approximately 80 onshore hours saved per week. The work included break investigation and trade-document scrutiny. This is a fund-level process metric, not a per-break time.

7. Level 3 - Agentic AI

Agentic AI changes the proposition from “follow this rule” to “investigate this exception using the available evidence, explain what happened, recommend the next action and escalate when confidence or materiality requires it.”

- IVP describes a hybrid model where a rules engine remains the system of record while specialised AI agents investigate and clear breaks, apply historical remarks and support human sign-off.
- Zensar's AWS Marketplace offering describes a multi-agent architecture that assembles break context, classifies root causes with confidence scoring, recommends corrective actions and routes high-impact or low-confidence cases to humans.
- The market already demonstrates commercial interest in agentic exception management. RESOLVIA should not claim to be first-to-market; its differentiation is the controlled comparative evaluation of agentic reasoning against conventional automation.

8. RESOLVIA Value Proposition

Value lever	Mechanism	Business benefit
Lower investigation effort	Agents gather records, documents, correspondence and precedents	Fewer analyst touches and more capacity
Faster resolution	Immediate triage and evidence retrieval	Higher same-day resolution; lower aged inventory
Consistency	Structured RCA and recommendation format	Less variation and stronger handovers
Knowledge reuse	RAG retrieves relevant historical cases	Recurring issues become reusable knowledge
Controlled automation	Risk-based approval gates	Automation without uncontrolled financial action
Verified closure	Expected vs actual post-action validation	Lower false closure and stronger audit confidence
Management visibility	Break, ageing, exposure, route and recurrence KPIs	Targeted automation and continuous improvement

9. Target Operating Model

Step	RESOLVIA	Human
Break created	Receive from existing reconciliation platform	-
Triage	Classify, prioritise, route	Override if needed
Investigation	Retrieve/connect evidence	Review material/ambiguous cases
RCA	Explain cause and alternatives	Challenge/approve reasoning
Resolution proposal	Recommend corrective action	Authorise consequential action
Execution	Invoke permitted controlled workflow	Maker-checker approval
Validation	Compare expected vs actual	Review exceptions/reopens
Closure	Close only when validation passes	Accountability remains with operations

Design principle: AI investigates. Humans authorize. Systems execute. AI validates.

10. Business Impact Model

The project will not invent a universal dollar saving. Once pilot measurements are available, the business case can be calculated as:

Annual capacity released = eligible break volume × average manual investigation time × automation coverage × working-day factor.

Potential cost avoidance = annual capacity released × loaded analyst cost per hour

Additional value can be modelled from reduced ageing, fewer escalations, lower operational risk and improved same-day resolution.

External figures establish that the opportunity exists; RESOLVIA's controlled measurements establish whether an agentic layer creates incremental value.

11. Pilot KPIs

KPI	Management relevance
Investigation Cycle Time	How quickly a break moves from detection to decision
Investigation Effort Reduction	Analyst capacity released
Straight-Through Investigation Rate	Share moving without manual investigation
Resolution Recommendation Accuracy	Quality alongside productivity
Break Classification Macro-F1	Reliability across break families
RCA Accuracy / Evidence Completeness	Correctness and supportability
Same-Day Resolution %	Operational deadline performance
Aged Break Rate	Queue reduction
False Closure / Reopen Rate	Closure quality
Unauthorized Action Rate	Must remain 0 for consequential actions
Audit-Trail Completeness	Evidence → recommendation → approval → action → validation trace

12. Why Agentic AI Rather Than More RPA?

- RPA is strongest when the path is known and repeatable.
- Real exception queues contain missing fields, contradictory documents, timing differences, multiple systems and ambiguous correspondence.
- Rules determine what to do when the scenario is known; agents investigate when the path must first be discovered from evidence.
- RESOLVIA does not replace RPA. It sits above deterministic automation and identifies when deeper investigation is required.
- The objective is maximum safe automation, not AI everywhere.

13. Pilot Proposal

1. Select one cash-equity exception family with sufficient historical variety.
2. Build a controlled synthetic operational environment mirroring structured and unstructured evidence.
3. Establish manual-style and deterministic rules/RPA-style benchmarks.
4. Deploy RESOLVIA with ML, OCR, RAG, agentic reasoning and human approval gates.
5. Run the same controlled cases through all three approaches.
6. Compare time, effort, coverage, RCA quality, recommendation accuracy, false closure and control outcomes.
7. Translate measured improvements into analyst capacity and cost impact.

14. Sources & Case Links

Clickable source list for the industry statistics and cases used in this pitch:

- [DTCC - Accelerated Settlement FAQs & Resources \(Q1 2026\)](#)
- [DTCC - Institutional Trade Processing Best Practices](#)
- [SS&C Blue Prism - Post-Trade Automation / SS&C GlobeOp case](#)
- [SS&C Blue Prism - SS&C GlobeOp trade reconciliation case](#)
- [CRISIL - Smart shadow accounting: 65% autonomous break resolution](#)

- [CRISIL - Simplified reconciliation: 55% cost savings](#)
- [CRISIL - 99% TAT reduction / ~80 hours saved per week](#)
- [CRISIL - Asset Servicing & Operations](#)
- [Broadridge - State of the Market: FCM Front Office](#)
- [Capco - Reconciliations: Five Trends Shaping the Future Landscape](#)
- [IVP - Reconciliation Solution](#)
- [AWS Marketplace / Zensar - Intelligent Trade Reconciliation](#)

15. Case Viability in One Page

Question	Business answer
Is the problem real?	Yes. Published cases show large break volumes and substantial investigation/operations effort.
Is it already automated?	Partly. Matching and repeatable resolution are mature in many environments, but maturity varies.
Why another layer?	The residual exception population is the least standardized and most evidence-heavy.
Why agentic AI?	It can dynamically retrieve, connect and reason across heterogeneous evidence.
Why not full autonomy?	Consequential financial/client actions require accountability; human approval remains mandatory.
How will value be proven?	Controlled comparison against manual-style and rules/RPA-style baselines.
Strategic upside?	A scalable exception-intelligence layer that can increase automation coverage while preserving governance.

RESOLVIA - From Break to Verified Resolution